PRODUCTBC AI BUILDATHON · MARCH 2026

Connect4

Neighbourhood Fit Tool

Onboarding Flow — Annotated Screen Walkthrough

Prototype v4 · 14-question flow · 4 phases

WHAT'S IN THIS DOCUMENT

- 15 annotated screens — cover, landing, 13 flow screens
- Faithful reproduction of all questions, options, and copy
- Screen intent, design rationale, and UX decisions per panel
- Build limitations and placeholder states clearly flagged

FLOW STRUCTURE

PHASE 1	Practicalities	Q1–Q5	Why coming · Timeline · Household · Budget · Transport
PHASE 2	Your Lifestyle	Q6–Q10	Free day · Neighbourhood energy · Outdoors · Cultural community · Comfort & safety
PHASE 3	Where you're from	Q11–Q12	Current location · Describe to a friend
PHASE 4	Your favourite place	Q13–Q14	Favourite location · What makes it yours → See my match

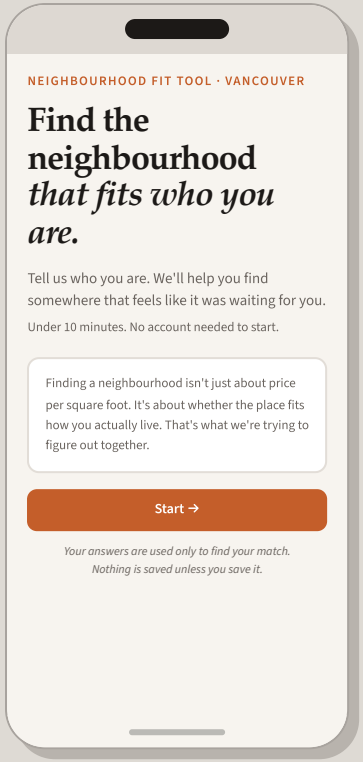

SCREEN INTENT

The landing screen has one job: earn permission to ask personal questions. It must establish trust, signal the product's value proposition, and remove friction from the decision to start — all before the user has answered a single question.

DESIGN **Serif headline, sans-serif body**

Palatino Linotype is used for the headline only — a system serif that loads instantly with no web font penalty. It signals warmth and permanence against the utilitarian body type. The italic on the second line breaks the formal register and adds emotional resonance without changing typefaces.

UX **Privacy note below the CTA, not above**

Placing the privacy note after the action button avoids leading with anxiety. Users who are ready to start don't need to wade through a disclaimer. Users who hesitate will scroll down and find the reassurance there. Trust is earned through behaviour, not upfront statements.

UX **Single CTA, no secondary actions**

The landing screen offers one choice: start. No login prompt, no "learn more," no navigation. Reducing decision load at the entry point directly improves start rates. Account creation is deferred entirely to the result screen's Save action.

What's bringing you to Vancouver?

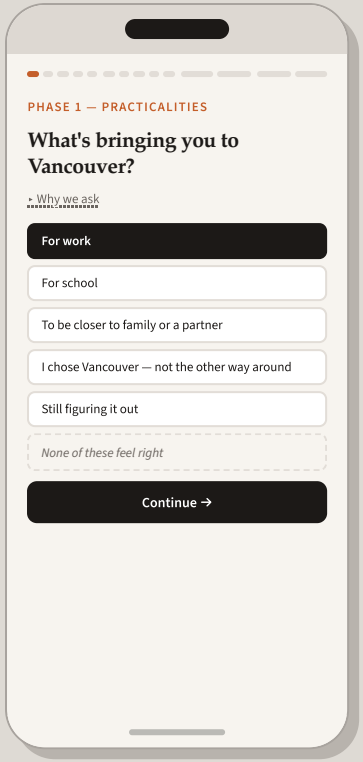

SCREEN INTENT

The first question establishes context for every result that follows. The reason for moving to Vancouver shapes whether proximity to community, transit reliability, or neighbourhood character matters most. This is the broadest filter in the matching logic.

DESIGN **Progress indicator — pip bar**

14 pips grouped into 4 phase clusters, separated by small gaps. Phase grouping communicates that the flow has structure and rhythm — not an undifferentiated list of 14 questions. Research consistently shows that visible progress reduces abandonment mid-flow.

UX **"Why we ask" inline disclosure**

A collapsible toggle appears on every question. Users who trust the product skip it; users who feel uncertain can read it. The disclosure is written in plain language and explains specifically what the answer changes — not generic reassurance. This is especially important for a vulnerable population making high-stakes decisions.

UX **Escape hatch on every multiple-choice question**

Every MC question includes a "None of these feel right" option that reveals an inline free-text field. No preset list can cover every newcomer's situation. The escape hatch prevents users from either abandoning the flow or choosing an inaccurate answer under pressure — both of which would degrade match quality.

Timeline + Household

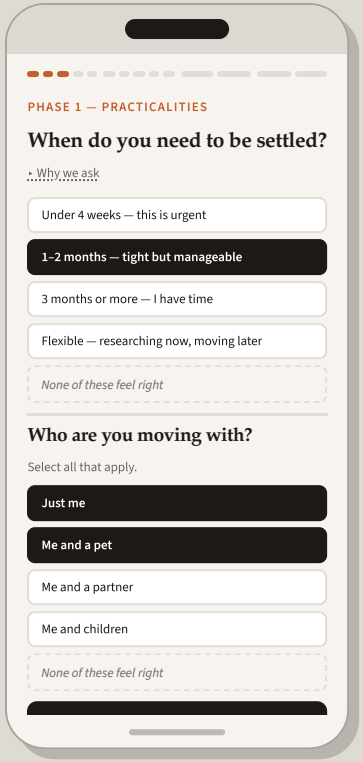

SCREEN INTENT

Timeline and household composition are practical constraints that modify the matching output — not lifestyle signals. A user with three weeks and a toddler needs a different result to one who is researching for six months alone. These two questions are grouped here to maintain pace through the practical phase.

TECH — Timeline as a result-weighting input

The timeline answer is not used to filter neighbourhoods — it adjusts how the matching logic weighs confidence vs. breadth. Short-timeline users receive higher-confidence, lower-variance recommendations. Long-timeline users receive a wider exploration set with more nuance. Same data, different output strategy.

UX — Multi-select for household

Household composition uses multi-select because real households are compound: a person moving alone with a dog is different from a person moving alone. Multi-select captures this without adding questions. The matching logic uses this to weight school proximity, park access, and noise tolerance accordingly.

DESIGN — Two questions, one screen

Phase 1 contains 5 practical questions. Pairing shorter questions on one screen where content allows maintains momentum and avoids the sense that the form is endless. Only done where questions are thematically grouped and neither requires significant reflection.

Budget + Transport

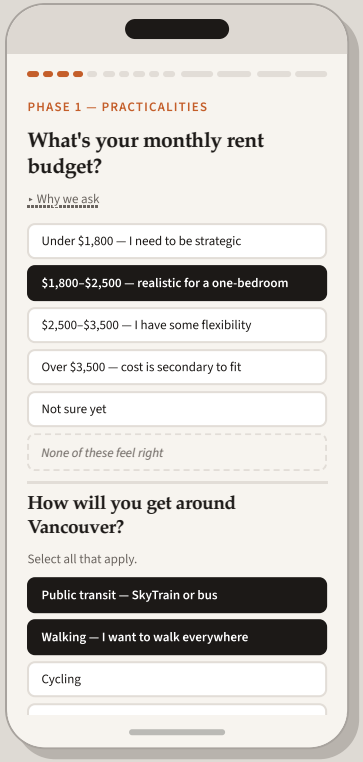

SCREEN INTENT

Budget and transport are hard geographic constraints. Together they define which neighbourhoods are physically accessible and financially viable — before any personality matching begins. They are the floor and walls of the result set.

DESIGN — Budget options carry emotional tone

Rather than displaying clinical dollar ranges alone, each option includes a plain-language description of what that budget means in the Vancouver market. *"I need to be strategic"* and *"cost is secondary to fit"* let users self-identify their financial situation without having to benchmark against local norms they don't yet know.

TECH — Transport as a geographic filter

Transport mode directly constrains neighbourhood viability. A user without a car cannot realistically live somewhere without frequent transit. The matching layer uses transport answers to eliminate or de-weight neighbourhoods before presenting results — not as a preference signal but as a hard accessibility constraint.

UX — No escape hatch on transport

"Not sure yet" is included as a standard option, making the escape hatch unnecessary. Transport is also less likely to produce edge cases than more personal questions. Removing the escape hatch here slightly reduces visual noise without sacrificing inclusivity.

Picture a free day

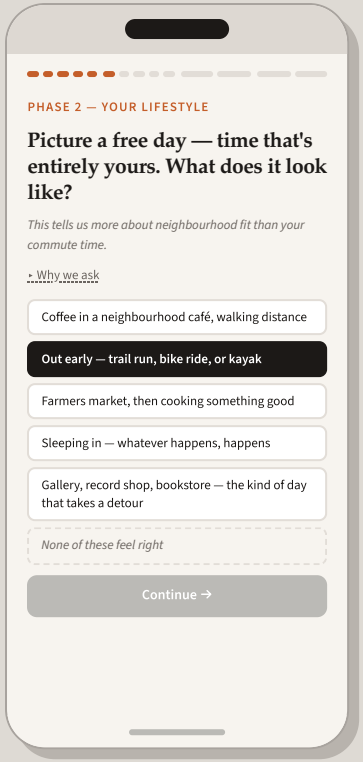

SCREEN INTENT

Phase 2 opens with a deliberate register shift — from logistics to identity. The free day question is the first in the flow designed to reveal how someone actually wants to live, not just what they need. It is one of the two highest-signal inputs in the matching logic.

UX "Free day" — inclusive framing

"A free day — time that's entirely yours" rather than "a Saturday." The Saturday framing excludes shift workers, users with religious observances, and parents with non-standard custody arrangements. Inclusive question framing matters particularly for an international newcomer population with diverse work patterns and cultural practices.

TECH Options map directly to neighbourhood personality types

Each answer option is weighted against a neighbourhood personality matrix. *"Out early — trail, bike, kayak"* maps to high outdoors access and lower density. *"Gallery, record shop, bookstore"* maps to cultural density and walkability. The options are not arbitrary — they were designed to produce discriminating signals across Vancouver's neighbourhood range.

DESIGN CTA shown disabled until selection is made

The Continue button renders in a muted grey until the user makes a selection. This communicates that an action is required without showing a validation error. No error message is needed — the disabled state is self-explanatory and avoids any sense of failure or correction.

Neighbourhood energy

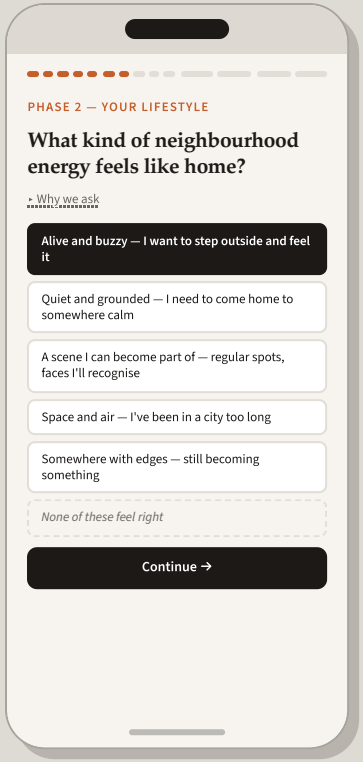

SCREEN INTENT

Neighbourhood energy is the core personality-matching question. It captures the felt quality of a neighbourhood that cannot be measured by any data point — density, noise, social tempo. Together with Q6 (free day), this question drives the largest share of the match output.

DESIGN — Options written as felt experiences

The options use evocative, first-person language rather than category labels. *"I've been in a city too long"* and *"still becoming something"* invite self-recognition rather than categorisation. This framing is more accurate and more engaging — users are more likely to select the option that actually matches their preference when they can picture themselves in it.

TECH — "Somewhere with edges" — emerging neighbourhood signal

This option specifically targets users open to neighbourhoods that are gentrifying or transitioning. From a matching standpoint it opens up a set of lower-cost, higher-upside areas that other options would not surface. It is also the least obvious option — users who choose it tend to have high tolerance for imperfect infrastructure in exchange for character and affordability.

UX — Single-select, not ranked

Users are asked to pick one energy type, not rank several. Ranking introduces cognitive load and often produces indecisive middle rankings that are analytically noisy. A forced single choice produces a cleaner signal, and users who genuinely can't choose are served by the escape hatch.

Outdoors access + Cultural community

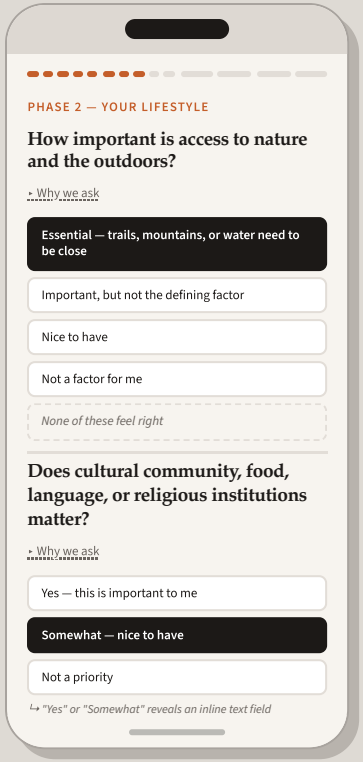

SCREEN INTENT

These two questions address Vancouver-specific geography and cultural belonging — factors that cannot be inferred from personality signals alone. They require direct questions because the answers are not predictable from any other input.

TECH Outdoors access reshapes geographic weighting

Vancouver is unusual in that mountains, ocean, and trails are genuinely accessible from many city neighbourhoods — but not all equally. A user who says outdoors access is essential will receive results weighted toward neighbourhoods with proximity to trails, beaches, or parks, even if those areas score lower on other lifestyle dimensions.

UX Conditional inline follow-up on Q9

Selecting *"Yes"* or *"Somewhat"* reveals an open text field inline — no new screen. The field captures specificity: a Mandarin-speaking community, a halal butcher, a Sikh gurdwara. This cannot be captured by preset options and requires free text. The conditional reveal avoids presenting the text field to users for whom the question isn't relevant.

UX Q9 framing — direct without apology

The question is asked plainly and confidently. The "Why we ask" text states: *"We ask because we can't reliably infer it from any other answer. For many newcomers, this question matters more than any other."* Hedging or softening this question would signal that it's somehow inappropriate — it isn't. For a large portion of the target user group, this is the most consequential question in the flow.

Comfort & safety

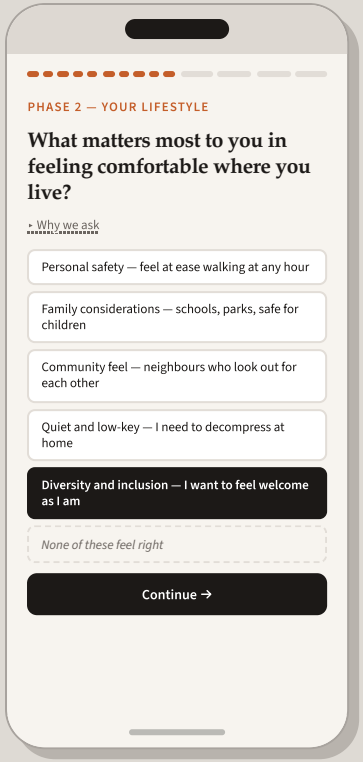

SCREEN INTENT

Q10 closes Phase 2 by addressing a dimension of fit that no other question captures: what "safe" and "comfortable" means to this specific person. Safety is not a single data point — it is a felt quality that means different things to different people.

DESIGN **"Comfortable" not "safe"**

The question stem uses *"comfortable"* rather than *"safe"* as the framing word. *"Safe"* implies threat — a loaded signal for a user who has not yet visited Vancouver and has no reference point for its neighbourhoods. *"Comfortable"* encompasses safety without anchoring to the anxiety framing. It also permits options like community feel and diversity, which are not safety questions at all.

UX **Diversity and inclusion as a first-class option**

For LGBTQ+ newcomers, visibly non-Western arrivals, or users who have experienced discrimination in previous homes, neighbourhood acceptance is the defining question. No other item in the flow captures this signal. Making it a first-class option — not a sub-question or opt-in — signals that the product was designed with these users in mind.

TECH **Phase 2 complete — all 10 lifestyle pips filled**

The pip bar now shows all 10 Phase 1 + Phase 2 pips as active. This visible milestone gives users a sense of accomplishment before the flow shifts to the more reflective open-text phases. The pip bar serves as a passive progress reward without any explicit congratulatory message.

Current neighbourhood — location

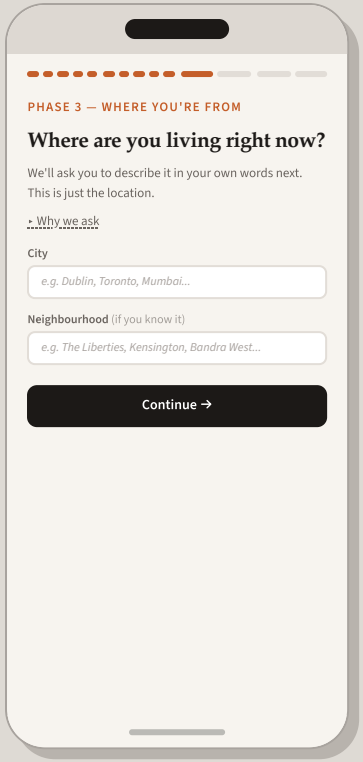

SCREEN INTENT

Phase 3 introduces the analogous neighbourhood concept — the product's highest-value differentiator. Before asking the user to describe their current neighbourhood, the location must be captured as structured data for geocoding and contextual lookup.

TECH Structured input for geocoding

City is a required field; neighbourhood is optional. City enables geographic context lookup. Neighbourhood, when provided, feeds directly into the result screen's analogous comparison: *"Think of it as a version of [their neighbourhood]…"* The note *"(if you know it)"* removes pressure for users who identify with a city but not a specific area.

UX Expectation-setting note

"We'll ask you to describe it in your own words next. This is just the location." This note appears inline above the fields. It prepares the user for Q12 (the open-text description) while signalling that this screen is low-effort — just type where you live. Reducing perceived complexity on any given screen improves completion.

DESIGN No Vancouver neighbourhood names in the flow

A core constraint of the entire onboarding: no Vancouver neighbourhood names appear anywhere before the result screen. The product is designed for users with zero local knowledge. Showing neighbourhood names mid-flow would create anxiety for users who don't recognise them and false confidence for those who half-recognise them.

Describe your neighbourhood to a friend

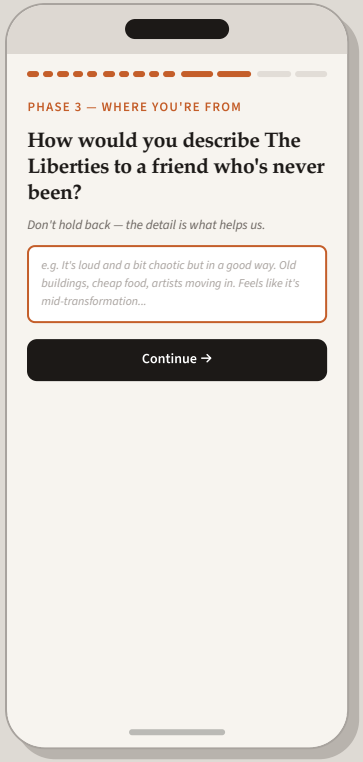

SCREEN INTENT

This is the first open-text question in the flow and the first moment where AI processing becomes the primary mechanism. The user's description of their current neighbourhood — in their own words — is the richest single input the product receives. No structured question could capture the same signal.

TECH — Free text → AI semantic interpretation

This input is passed to the AI layer at result time, not in real-time. The AI interprets the qualitative description — density, energy, social character, pace — and maps it against a neighbourhood personality matrix. A description like *"loud, chaotic, artists moving in"* produces a very different weighting than *"quiet streets, families, everyone knows each other."*

DESIGN — Placeholder as modelled behaviour

The placeholder text is written in the voice of an actual answer, not as a field label. This models the depth and tone the product is looking for. Users who are uncertain what to write can use the placeholder as a template. It also signals that imperfect, unpolished language is welcome — *"a bit chaotic but in a good way"* is the right register.

LIMIT — Real-time AI response not yet active

In a future version of the product, the AI will respond to the user's description with a brief reflection — appearing as an interstitial screen above the next question. This creates the conversational feel of a knowledgeable local listening and reflecting back. The current build defers all AI interpretation to the result screen.

Favourite neighbourhood — location

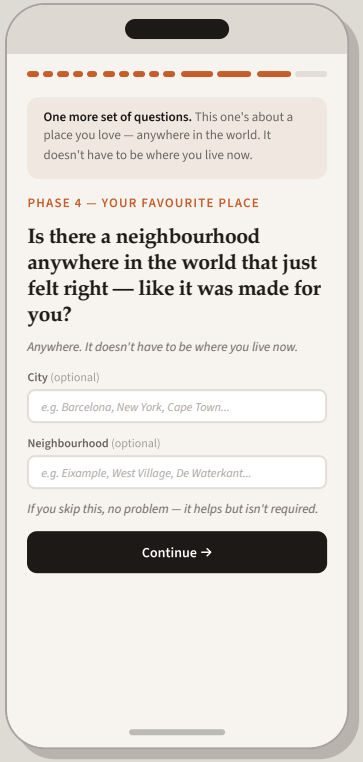

SCREEN INTENT

Phase 4 captures aspirational preference — what the user actually wants, unconstrained by their current situation. The favourite neighbourhood question reveals the emotional register they are hoping to find in Vancouver, and is the primary input for the analogous neighbourhood comparison on the result screen.

DESIGN **Bridge card signals phase transition**

A warm-toned card in the rust-light palette appears at the top of the screen to signal a shift in register — from "where you are" to "where you love." No interstitial screen is needed; the card does this work inline. It also reduces anxiety by clarifying that the next questions are reflective and exploratory, not prescriptive.

UX **Both fields optional, skip note present**

Not every user will have a meaningful favourite neighbourhood to name. Marking both fields optional and adding a skip note removes the pressure to perform. Users who do answer provide a richer signal; users who skip don't feel penalised. The matching logic handles absent data gracefully by weighting other inputs more heavily.

TECH **Location feeds the analogous comparison**

The city and neighbourhood entered here, combined with the Q14 description, generates the analogous neighbourhood comparison block on the result screen. The template — *"Think of it as a version of [their place] — but with [X] and [Y]"* — requires this structured location data to produce a meaningful, personalised comparison.

What makes it your favourite → See my match

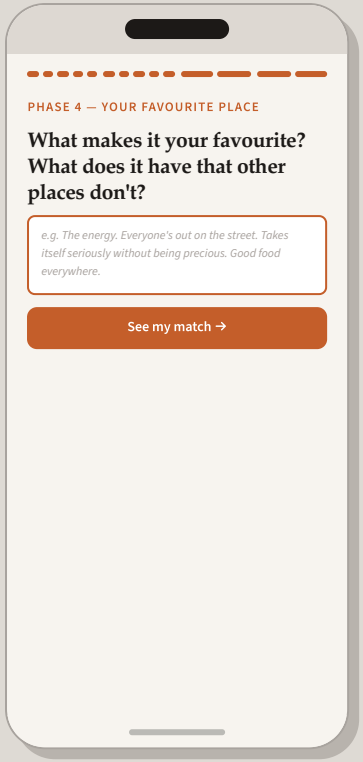

SCREEN INTENT

The final question in the flow captures the user's aspirational standard — the qualities they value most in a neighbourhood, expressed in their own words and without constraints. All 14 pips are now filled, and the CTA fires directly to the loading screen. This is the emotional peak of the onboarding.

- TECH **Highest-signal open text input**

  This description, combined with Q12 (current neighbourhood), forms the core of the AI's qualitative matching. The user is describing what they love about a place with no constraints — this reveals their deepest preferences more accurately than any structured question. The AI maps these qualities against neighbourhood personality profiles in the data layer.

- DESIGN **CTA label changes — "See my match →"**

  The final CTA uses different language to the previous Continue buttons. *"See my match →"* signals that this is the end of input and the beginning of output. The rust colour (used elsewhere only for primary actions) reinforces the significance. This is the moment the user has been building toward — the CTA should feel like a reveal, not just another step.

- UX **No confirmation screen**

  Pressing *"See my match →"* fires directly to the loading screen — no intermediate confirmation step. At this point the user has invested 8–10 minutes and completed all 14 pips. Interrupting with a confirmation screen would deflate the momentum built through the flow. The fully-filled pip bar is confirmation enough.

Loading / Processing Screen

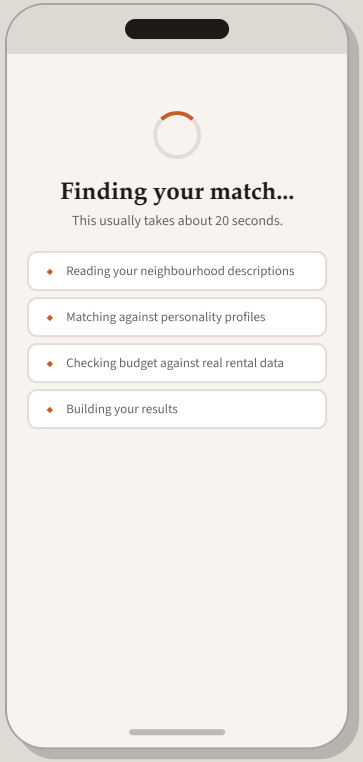

SCREEN INTENT

The loading screen bridges the end of onboarding and the delivery of results. Its job is to maintain trust and manage the wait — AI processing and data layer queries take real time. Rather than leaving the user with a generic spinner, the screen makes the process legible.

UX **Named processing steps build transparency**

Four chips animate in sequentially, each naming a specific operation: reading descriptions, matching profiles, checking rental data, building results. Naming the steps communicates that something meaningful is happening — not a simulated delay. For users making a high-stakes decision, visibility into the process is a trust signal in itself.

DESIGN **Staggered chip animation**

Chips appear with staggered fade-in delays (0s, 0.5s, 1.0s, 1.5s). This creates the visual impression of real-time computation — each step appearing as it completes. The effect is not deceptive; it is a design representation of a process that is genuinely sequential. The animation also gives the user something to watch during the wait.

UX **No cancel option**

No back button or cancel action is offered. At this point the user has completed the full flow and is waiting for output. Offering a cancel option introduces doubt at exactly the wrong moment. If the process fails, an error state handles recovery — but under normal conditions, the loading screen is a one-way passage to results.

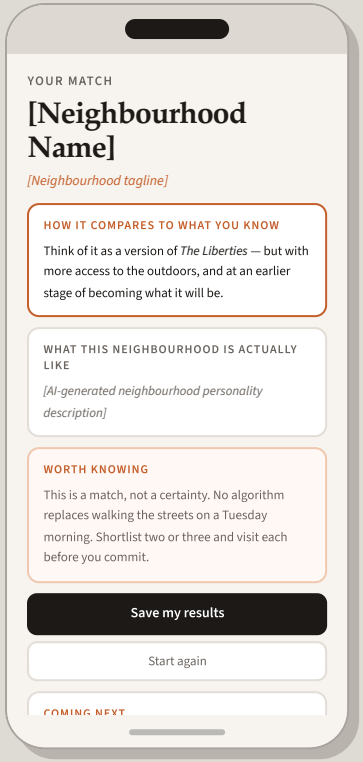

SCREEN INTENT

The result screen delivers the product's core promise: a personalised, qualitative neighbourhood match. It must feel earned — substantive enough to justify the 14-question investment — while being honest about what a match can and cannot guarantee. **The entire content layer of this screen is a placeholder in the current build.** All text blocks are populated by the AI at match time in production.

TECH **Analogous neighbourhood — the core differentiator**

The rust-bordered block is the product's primary output. The template — *"Think of it as a version of [their neighbourhood] — but with [X] and [Y]"* — produces a comparison that no existing tool can generate. It references the user's own lived experience and translates the match into terms they already understand. In production this is AI-generated from Q11–Q14 inputs.

UX **"Worth knowing" — honesty as trust**

The caveat block explicitly states the product's limitations. *"This is a match, not a certainty"* and the recommendation to visit in person before committing are deliberate. A product that acknowledges its limits is more trustworthy than one that overpromises. For users making high-stakes housing decisions, this is a feature, not a disclaimer.

LIMIT **Entire content layer is placeholder**

In the current prototype build: neighbourhood name, tagline, analogous comparison text, and the personality description block are all static placeholders. In production, all four are generated by the AI layer at match time, using the user's full input set against the neighbourhood data layer. This screen tests layout, hierarchy, and emotional arc — not content.

UX **Account gate at Save, not before**

No account is required to read the full result. The user receives complete value before being asked to create an account — the prompt appears only when they press Save. This earns trust before requesting commitment, and keeps the result screen free of friction for the majority of users who won't save on a first visit.